
How using a Graph Convolutional Network can actually enhance the Quantum Approximation for Ising Models.

Rebecca Tsekanovskiy, Micheal Halpern
Rensselaer Polytechnic Institute
Troy, NY
Tsekar@rpi.edu, Halpem2@rpi.edu

Abstract

1 The Ising Spin Model is a pivotal framework for understanding the complex interac-
2 tions of particles in spin glass systems, with applications spanning physics, biology,
3 and even economics. In our research, we leverage the Quantum Approximation
4 Optimization Algorithm. To solve the Ising Model by approximating the minimiza-
5 tion of its Hamiltonian, thereby efficiently finding the system's ground state. A key
6 challenge in optimizing QAOA is determining the optimal γ and β , which govern
7 the algorithm's performance. Traditional methods for selecting these parameters
8 are computationally intensive and often impractical for large-scale systems. To
9 address this, we implemented a graph convolutional network to predict optimal γ
10 and β values to analyze the expected value for a max cut based off the predictions.
11 Ultimately, by utilizing a graph convolutional network with the algorithm, we were
12 able to improve the calculation time of γ and β values by 99.9% compared to a
13 classical approach, and a 1% increase in expected value. These results demonstrate
14 the potential of machine learning-enhanced quantum algorithms to tackle large,
15 complex optimization problems with unprecedented efficiency.

16 1 Background

17 1.1 Ground State of the Ising Model

18 The Ising Spin Model is a critical component in the study of Spin Glass systems, which are concerned
19 with the interactions of particles. In this model, each spin, denoted as σ_i , can take one of two values:
20 +1 or -1, where σ_i represents the i-th spin. Particles with a spin quantum number of -1 have a spin
21 orientation that is opposite to that of particles with a spin quantum number of +1 [5]. The term w_{ij}
22 signifies the interaction between neighboring spins and their strength. It holds a value of zero for
23 non-neighboring spins and -1 if and only if spins i and j are connected [5]. Each spin σ_i contributes
24 its energy to the Hamiltonian of the spin system. This contribution is defined as follows [3]:

$$H(\sigma) = - \sum_{i < j} w_{ij} \sigma_i \sigma_j \quad (1)$$

25 This equation represents the Hamiltonian of the system, which is a measure of the total energy of the
26 system. The ground state of the Ising model corresponds to the configuration of spins that minimizes
27 this Hamiltonian and thus minimizes the total energy of the system. Finding the ground state of an
28 Ising spin model can be mapped to a two-dimensional graph, where each particle can be represented
29 as a node. Particles with spin quantum numbers of +1 and -1, which are neighbors, share an edge in
30 the graph. Minimizing the energy of the physical system and finding its ground state is equivalent to
31 finding the max-cut of this graph. An example of mapping the Ising spin model to a two-dimensional
32 graph and finding the graph's max-cut is 1.




Figure 1: Example: A graph with five nodes and six edges. The dashed line represents the max-cut of this graph: the empty nodes are in set V_1 and the colored nodes are in set V_2 . There are other max-cuts, but in this graph the max-cut is five.

1.2 Quantum Approximation Optimization Algorithm

The Quantum Approximation Optimization Algorithm (QAOA) is used for many different types of optimization problems in Quantum Computing, more specifically, the QAOA can generate an approximate Hamiltonian minimization and thus in turn approximately minimize the system's energy and find the approximate ground state of the Ising Model [5]. This is represented as 1 where $\mathbf{b} \in \{0, 1\}^n$ [5].

$$C(\mathbf{b}) = \sum_{\alpha=1}^m C_{\alpha}(\mathbf{b}) \quad (2)$$

Equations 1 and 2 are closely related, as 1 represents the Hamiltonian that 2 is approximately minimizing. By mapping the Ising system model to the max-cut problem, we can define the QAOA by the following steps [1]:

1. First, a qubit is allocated for each node in the graph. All qubits are initialized to a state that is an equal superposition.
2. Then, by applying the $U(H_c, \gamma)$, otherwise known as the Hamiltonian cost 3, we are able to calculate the cut value for the given graph for an angle γ .
3. By applying the $U(H_c, \beta)$, otherwise known as the Hamiltonian mixer 4, we are able to modify the current states of the qubits based on the cost calculated in step 3 for an angle β .
4. We will repeat the previous two steps, for a total of p times with different parameters γ_i and β_i , ultimately, creating new states. p is defined as the depth chosen for the algorithm.
5. Based off the function of the depth, $F_p(\gamma, \beta)$, we calculate the expected Hamiltonian cost after running the QAOA algorithm for a given depth p .

1.2.1 Hamiltonian Cost Function and Hamiltonian Mixer Function

The Hamiltonian Cost Function 3 acts on qubits i, j which are the edges in our graph and w_{ij} . This function applies Pauli-z operators on both nodes i and j rotating the qubits about the Z-axis for a given rotation γ . This function changes the qubits states based on their edge connections to other qubits [3].

$$H_C = \frac{1}{2} \sum_{\langle i, j \rangle} w_{ij} \left(I - \sigma_z^{(i)} \sigma_z^{(j)} \right) \quad (3)$$

The Hamiltonian mixer function 4 acts on all qubits in the circuit. A Pauli-X gate is applied, rotating the qubits about the X-axis for a given rotation β . This rotation results in mixing the computational state, changing the probability of receiving a positive or negative state when measuring the qubit [3].

$$H_B = \sum_{j=1}^n \sigma_x^{(j)} \quad (4)$$




Figure 2: The visual representation of the γ and β value search space for a given graph. In the example, 25 beta and 25 gamma values are calculated making the size of the search space 2500. The darker regions represent γ and β values associated with lower energy states and therefore more promising γ and β values, while the lighter regions represent γ and β values associated with higher energy states [7].

60 1.2.2 Finding Gamma and Beta Values

61 To find suitable γ and β values, we can search the two-dimensional solution space of those values.
62 To search the solution space, we can repeatedly calculate the expected value produced by running
63 the QAOA algorithm at depth one for a given set of γ and β values 2. Searching the solution space
64 this way allows us to find γ and β values which produce low amounts of energy after running QAOA
65 at depth one. These low energy values are likely to lead to lower states of energy when the QAOA
66 algorithm is run at a higher depth, producing a low energy state [3].

67 1.3 Graph Convolutional Networks

68 Graph Convolutional Networks (GCNs), a type of deep learning model operates specifically in a
69 non-Euclidean space, meaning these models best understand data in a graph like structure, like trans-
70 portation and social networks [8]. GCNs cause each node to continuously update its representation
71 based off the neighboring nodes around. This iterative process makes a GCN an effective choice for
72 capturing complex relationships in Ising spin models. We use GCNs to help find the most optimal γ
73 and β values for the QAOA algorithm. By using a GCN model, hidden patterns in the graph structure
74 are better understood and analyzed [9].

75 2 Related Works

76 This section discusses two studies focused on utilizing graph neural networks (GNNs) with the QAOA
77 algorithm.

78 2.1 Graph Neural Network Initialization of Quantum Approximate Optimization Algorithm

79 In this study, Jain et al. propose an innovative approach to initializing QAOA for solving the
80 Max-Cut problem using GNNs. The research addresses two key challenges: how to effectively
81 initialize QAOA and how to optimize its parameters to achieve the best possible solution [2]. The
82 authors demonstrate that integrating GNNs as a warm-start technique significantly enhances QAOA's
83 performance compared to using QAOA or GNNs independently. Additionally, they show that GNNs
84 provide warm-start generalization across various graph instances and sizes, an advantage not readily
85 available with other methods [2].

86 The methodology proposed involves using GNNs to initialize the QAOA's parameter space close
87 to an optimal solution, thereby reducing the quantum resources needed for optimization [2]. The
88 authors compare their GNN-based initialization method to other approaches, such as those using
89 semi-definite programming (SDP) relaxations and Trotterized quantum annealing. Through extensive
90 numerical experiments, they demonstrate that the GNN-augmented QAOA consistently outperforms
91 these traditional methods, particularly in scenarios where the depth of the quantum circuit is limited,
92 a common constraint in noisy intermediate-scale quantum (NISQ) devices [2]. Further, the study

highlights the scalability and generalization capabilities of GNNs when used for QAOA initialization. Unlike other methods that may require new problem instances for each new graph, the GNN approach generalizes well across different graph sizes, making it a versatile tool for various combinatorial optimization problems [2]. The research shows that even when trained on smaller graphs, the GNNs can effectively warm-start QAOA for larger and more complex graphs, maintaining high solution quality with reduced computational overhead [2].

In summary, Jain et al.’s work provides a compelling case for the integration of classical machine learning techniques, such as GNNs, with quantum algorithms like QAOA [2]. Their findings suggest that this hybrid approach can significantly enhance the efficiency and scalability of quantum optimization algorithms, particularly in solving NP-hard problems like Max-Cut. The study paves the way for future research into the combination of classical and quantum computing techniques to tackle complex optimization challenges.

2.2 Graph Learning for Parameter Prediction of Quantum Approximate Optimization Algorithm

The primary objective of this work is to enhance QAOA’s performance by optimizing its initialization process, which is crucial for efficient execution on NISQ devices. These devices, while offering a promising platform for quantum computations, are limited by their relatively small number of qubits and susceptibility to noise. By leveraging the strengths of GNNs, the researchers aim to provide a well-informed starting point for the QAOA parameters, thereby reducing the quantum computational overhead and improving the algorithm’s convergence rates.

The researchers focus on training four distinct GNN architectures—Graph Isomorphism Networks (GIN), Graph Convolutional Networks, Graph Attention Networks (GAT), and GraphSAGE, each selected for their unique strengths in capturing and processing graph-structured data. First, a synthetic dataset is created varying from sizes of 2 to 15. The methodology begins with data preprocessing, where the node degrees and one-hot encoding of node IDs are computed to serve as node features [4]. Once the dataset is prepared, it is fed into the various GNN architectures. Each GNN model is tasked with learning to predict the QAOA parameters, γ and β [4]. The GNNs perform inference by iteratively aggregating feature information from each node’s local neighborhood. This involves a message-passing framework where, at each layer, a node’s representation is updated based on the aggregated information from its neighbors.

For example, in the GCN model, the node representations are updated by applying a combination of linear transformations followed by a non-linear activation function. In contrast, GATs introduce an attention mechanism that assigns different weights to different nodes in the neighborhood, allowing the model to focus on the most relevant nodes during the aggregation process. The GraphSAGE model, on the other hand, optimizes for neighborhood sampling, allowing the GNN to efficiently handle graphs with large or varying node degrees [4]. After training, the GNNs output the initial estimates for γ and β , which are then used to warm-start the QAOA. This means that the QAOA does not begin its search from a random initialization but from a well-informed starting point, thus potentially avoiding poor local minima and improving the overall optimization process [4].

Liang et al. demonstrate that their GNN-based initialization not only reduces the computational burden on quantum devices but also exhibits strong generalization capabilities, successfully applying the model to larger, unseen graphs [4]. This scalability is particularly important for practical applications of quantum algorithms on NISQ devices.

3 Motivation

Quantum computing deals with optimizing problems like Ising Model, a known NP-complete problem. Solving the Ising model enables the solution of all other NP-complete problems [6]. A solution to The Ising Model can be found by using QAOA algorithms in linear time, but finding optimal β and γ parameters requires sampling a solution space. This research aims to optimize solving the Ising Model on the Rensselaer Polytechnic Institute (RPI)-IBM Quantum System One, to understand the challenges and limitations, and use machine learning algorithms to increase the efficiency of finding β and γ parameters for a given graph.

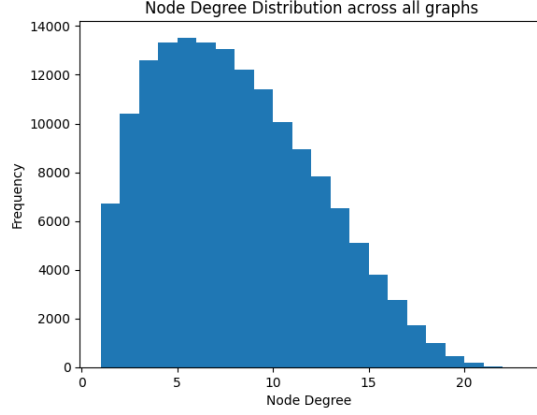


Figure 3: The histogram demonstrates the diversity of node degrees resulting from the predefined minimum and maximum fill rates and vertex counts.

Calculating the starting Gamma and Beta values usually require running the QAOA algorithm at depth one for a range of β and γ parameters. This process either requires a large scale simulation, which is very computationally expensive on classical hardware. This simulation scales linearly with the number of edges and nodes for a given graph. On the other hand, running on real quantum hardware is computationally very slow and expensive due to the repeated transpilation and long wait times for API calls, which are both required to search the β and γ parameter space. Both options are computationally expensive and slow when compared to a neural network, which should give results in a faster time.

4 Methodology

4.1 Data Creation

We created our own dataset containing a graph ID, adjacency list, and optimal γ and β values, which were used to train the graph convolutional network model. To generate this data, we defined four input parameters: minimum number of vertices, maximum number of vertices, minimum fill rate of a graph, and maximum fill rate of a graph. We define the minimum and maximum fill rates as the percentages of edges that are filled in when compared to a complete graph. A fill rate of 1 represents a complete graph. We selected the maximum number of nodes to be 24, as this is the limit of the Qiskit simulator, and the minimum number of nodes to be 2, as a graph requires at least two nodes to have a max cut greater than 0. To represent a broad spectrum of graphs, we set the minimum fill rate to 0.1 and the maximum fill rate to 0.9. This range of fill rates strikes a balance between sparsity and density in graph structures, enabling the generation of a diverse set of graphs with varying complexities. As shown in figure 3, the dataset generated includes graphs with varying densities and structures, as evidenced by the wide range of node degrees.

The number of vertices, defined as V , was randomly selected between the predefined minimum and maximum vertices. We then calculated the number of edges for V vertices. We set the minimum number of edges to be equal to the minimum fill rate times the number of edges in a complete graph. Similarly, we set the maximum number of edges to be equal to the maximum fill rate times the number of edges in a complete graph. We then chose a random number between the minimum and maximum number of edges, which is how the adjacency list was formed.

To calculate the optimal gamma and beta values, we ran the QAOA at depth one. We used COBYLA, a SciPy method for constrained optimization, to find the best gamma and beta values that provide a good expected value from the QAOA function. This computation was run on RPI's Aimos's DSC Cluster.

176 4.2 Model Architecture

177 The model is a Graph Convolutional Network implemented using PyTorch to predict optimal γ
 178 and β values for graph structures. The architecture is designed to capture complex relationships
 179 between nodes in a graph by leveraging multiple graph convolutional layers, each followed by batch
 180 normalization and ReLU activation functions to stabilize training and improve generalization.

181 We begin by initializing every node in the graph to have one singular feature, as we were more
 182 focused on how the structure of the graph rather than the individual node characteristics. The model
 183 consists of four graph convolutional layers, which are used to learn node representations based on the
 184 neighbors in the graph. The first convolutional layer looks at each node in the graph and all other
 185 neighbors of that node. By combining all the features it is now learned about the neighbors of the
 186 original node, the node then contains 128 features compared to the original one. By looking at the
 187 previous layer, this process repeats ultimately for the next four convolutional layers, going from 128
 188 features to 256 and finally outputting 512 features per node. By the time all the convolutional layers
 189 have been passed through, each node has gone from being represented as just a singular number to be
 190 represented by 512 feature vector.

191 The next part of the process is going through global pooling layer, which ultimately causes each
 192 nodes feature vector to be condensed into a single vector that summarizes the entire graph. As such,
 193 for each of the 512 features in the vector, the average is computed across all the nodes in the graph,
 194 ultimately creating a vector that contains the overall structure and feature distribution of the graph.
 195 More so, it reduces the dimensionality of the node features to a fixed size, regardless of the number
 196 of nodes in the graph. This layer is crucial for reducing the computational complexity. This layer will
 197 then output a singular 512 dimensional vector that represents the entire graph.

198 Lastly, the 512-dimensional vector passes through a series of fully connected layers, where it is first
 199 reduced to 256 dimensions. A ReLU activation is applied, allowing the model to capture complex
 200 patterns that might not have been recognized yet. This process repeats, with the vector being reduced
 201 to 128 dimensions, followed by another ReLU activation, and so forth until the vector is finally
 202 reduced to a size of two. These two dimensions correspond to the predicted γ and β values. The
 203 non-linear ReLU function is crucial for the model structure because most real-world relationships are
 204 non-linear. Without this, the model would only be able to capture linear relationships, which could
 205 lead to underfitting and limit its ability to predict complex relationships.

```

206
207 1 class GCNRegressor(nn.Module):
208 2     def __init__(self):
209 3         super(GCNRegressor, self).__init__()
210 4         self.conv1 = GCNConv(in_channels=1, out_channels=128)
211 5         self.conv2 = GCNConv(in_channels=128, out_channels=256)
212 6         self.conv3 = GCNConv(in_channels=256, out_channels=256)
213 7         self.conv4 = GCNConv(in_channels=256, out_channels=512)
214 8         self.bn1 = nn.BatchNorm1d(128)
215 9         self.bn2 = nn.BatchNorm1d(256)
216 10        self.bn3 = nn.BatchNorm1d(256)
217 11        self.bn4 = nn.BatchNorm1d(512)
218 12        self.fc1 = nn.Linear(512, 256)
219 13        self.fc2 = nn.Linear(256, 128)
220 14        self.fc3 = nn.Linear(128, 2) # Predicting gamma and beta
221 15
222 16    def forward(self, batch):
223 17        adj_matrices, params = batch
224 18        batch_size, num_nodes, _ = adj_matrices.size()
225 19
226 20        outputs = []
227 21        for i in range(batch_size):
228 22            adj_matrix = adj_matrices[i]
229 23            edge_index = adj_to_edge_index(adj_matrix)
230 24            node_features = torch.ones((num_nodes, 1)).to(device)
231 25
232 26            x = F.relu(self.bn1(self.conv1(node_features,
233            edge_index)))
234 27            x = F.relu(self.bn2(self.conv2(x, edge_index)))
235 28            x = F.relu(self.bn3(self.conv3(x, edge_index)))

```

```

23629         x = F.relu(self.bn4(self.conv4(x, edge_index)))
23730         x = F.dropout(x, p=0.5, training=self.training)
23831         x = F.relu(self.fc1(x))
23932         x = F.relu(self.fc2(x))
24033         out = self.fc3(x.mean(dim=0))
24134         outputs.append(out)
24235     return torch.stack(outputs)
243

```

244 4.3 Implementation

245 The model training was implemented in PyTorch, focusing on ensuring compatibility with RPI's
246 AiMos NPL Cluster and the associated graphics processing unit (GPU). An instance of the GCN
247 model was instantiated and moved to the correct GPU device. Using a learning rate of 0.01 and the
248 Adam optimizer, we began to train our model. We had a total of 20 epochs and a batch size of 32.
249 For each batch, the model would make its predictions, and the loss was calculated using a custom
250 loss function specific to the batch. We trained the model using a single node on the RPI Center for
251 Computational Innovation's NPL Cluster. This node was equipped with two 20-core 2.5 GHz Intel
252 Xeon Gold 6248 processors, eight NVIDIA Tesla V100 GPUs (each with 32 GiB HBM), and 768
253 GiB of RAM. However, our PyTorch model only utilized one GPU for training.

254 4.3.1 QAOA Qiskit

255 Later in our code implementation, after finding optimal γ and β values, we generate a circuit that can
256 be run on the RPI-IBM Quantum System One. In this context, the GCN model predicts the γ and β
257 parameters, which are then used by the QAOA circuit function to construct, optimize, and run the
258 quantum circuit. Our circuit generation is defined as follows:

```

259
2601 def QAOA_circuit(beta: float, gamma: float, depth: int,
261 adjacencyList: list[list[int]]):
2622     num_qubits = len(adjacencyList)
2633     qc = QuantumCircuit(num_qubits, num_qubits)
2644     qc.h(range(num_qubits))
2655     qc.barrier()
2666     for _ in range(depth):
2677         for edgeIndex, neighbors in enumerate(adjacencyList):
2688             for edge in neighbors:
2699                 if edge > edgeIndex:
2700                     qc.cz(edgeIndex, edge)
2711                     qc.rz(2 * gamma, edgeIndex)
2722                     qc.cz(edge, edgeIndex)
2733     qc.barrier()
2744     for qubit in range(num_qubits):
2755         qc.rx(2 * beta, qubit)
2766     qc.barrier()
2777     qc.measure(range(num_qubits), range(num_qubits))
2788     return qc

```

```

280
2811 def QAOA(beta: float, gamma: float, fakeBackend: bool, depth: int,
282 adjacencyList: list[list[int]],) -> dict:
2832     if fakeBackend:
2843         backend = GenericBackendV2(len(adjacencyList))
2854     else:
2865         backend = renslearBackend
2876     qc = QAOA_circuit(beta, gamma, depth, adjacencyList)
2887     job = backend.run(qc, shots=1024)
2898     result = job.result()
2909     return result.get_counts()
291

```

Table 1: Comparison of Metrics for Graphs

Graph	Classic β	Classic γ	Classic EV	Classic Time(s)	GCN β	GCN γ	GCN EV	GCN Time(s)
Graph 1	0.6722	0.5703	28.42	19.45	-0.0655	0.1148	27.98	0.0038
Graph 2	0.7546	0.8523	39.09	47.93	-0.0340	0.0758	38.14	0.0030
Graph 3	2.2413	2.1869	11.87	17.36	-0.0320	0.1051	12.04	0.0029
Graph 4	0.6112	0.6382	17.14	20.15	-0.0630	0.1202	17.20	0.0027
Graph 5	3.9368	1.1151	8.86	8.82	-0.0235	0.1159	8.94	0.0027
Graph 6	2.4383	1.7958	57.18	58.56	-0.1189	0.2139	56.29	0.0029
Graph 7	1.5488	0.5639	33.91	260.22	-0.0436	0.1134	33.60	0.0027
Graph 8	1.4853	1.2505	64.99	62.35	-0.1073	0.1549	64.77	0.0027
Graph 9	0.4996	0.5010	15.05	11.77	-0.0376	0.1118	14.80	0.0027
Graph 10	0.7479	0.4673	65.07	347.86	-0.1058	0.1216	63.71	0.0028

5 Preliminary Results

Based off the total of ten graphs as shown in 1, which were all run on the RPI-IBM Quantum System One, the classical expected value (EV) was 341.5713, and the GCN EV was 337.4512. The total time to find the γ and β values using the classical approach was 854.4607 seconds, while the total time to find the γ and β with the GCN was 0.029002 seconds.

$$\text{Reduction} = \left(\frac{854.4607 - 0.029002}{854.4607} \right) \times 100 \approx 99.997\%$$

$$\text{Expected Value} = \left(\frac{341.5713 - 337.4512}{341.5713} \right) \times 100 \approx 1.2\%$$

Therefore, we achieved a 99.99% reduction time with a 1.2% expected value reduction. This suggests that the GCN method is not only computationally efficient but also maintains a high degree of accuracy, making it a viable alternative to classical approaches in practical scenarios. Although the GCN model demonstrates a significant reduction in prediction time, it is important to consider the training phase, which required approximately 521 minutes. This training time is a one-time cost, and once the model is trained, the benefits of rapid prediction times are realized across all future predictions. This trade-off between initial training time and subsequent prediction efficiency makes the GCN model a highly practical choice for large-scale applications.

6 Final Results

Method	Total Time (s)	Average Time (s)
GCN	1.4501	0.0290
Classical	42723.0359	854.4607
Random	0.0009549	0.0000191

Table 2: Comparison of total execution time and average execution time per graph between GCN, classical, and random methods over 50 randomly chosen graphs. The GCN approach significantly reduces the computational time compared to the classical method.

In our next round of testing, as shown in table 2 on average the GCN parameter method is approximately 1.0002 times better than classical methods and 1.003 times better than the random approach 5. Meanwhile, the GCN reduced the computation time required to determine optimal γ and β values by 99.996% across the 50 graphs, taking a classical method 42723.03585 seconds to 1.450097561 seconds to determine these values. The slight reduction in the expected value is offset by the considerable gains in computational efficiency, making a GCN model a compelling choice for large scale quantum computations. Ultimately, our testing demonstrates that through the usage of a GCN model, performance of total EV only improves and consistently outperforms other methods.




Figure 4: Comparison of the average expected values obtained using GCN-generated parameters, classically optimized parameters, and randomly selected parameters over 50 different graphs. The GCN approach shows slightly higher performance compared to the classical method, while random selection yields the lowest average expected value.




Figure 5: Total expected value comparison over 50 different graphs for GCN-generated parameters, classically optimized parameters, and randomly selected parameters. The GCN approach achieves the highest total expected value, slightly outperforming the classical method, while the random selection results in a significantly lower total expected value. The GCN has a total of 1734.856 expected value; Classical has a total of 1734.497 expected value across the 50 graphs. Random has a total of 1729.697 expected value.

7 Conclusion

In this research, we successfully leveraged a Graph Convolutional Network to predict optimal β and γ for the QAOA algorithm applied to the Ising Spin Model. Our approach demonstrated a significant reduction in computation time, achieving a 99.99% decrease compared to classical methods, with a 1% increase in expected value.

The solving of the Ising Spin System will map over to any other NP-complete problem. Ultimately, by improving the computation of this problem will help the scientific community by improving the solutions to problems like polymer folding, memory, and decision-making in social sciences and economics[6].

In conclusion, our work not only advances the field of quantum optimization but also holds promise for broader scientific applications. Our work with QAOA and GCNs aims to develop a model that not only can calculate β and γ values much faster than traditional methods, but also serves as a benchmark for future research in quantum optimization.

8 Future Work

8.1 Monte Carlo k-folding

In our current approach, the custom loss function operates by randomly selecting 5 points from the dataset to evaluate the model's predictions. While this method provides a quick and straightforward way to approximate the model's performance, it has limitations. Specifically, because only a small, random subset of points is considered in each iteration, there's a risk that the model might not fully capture the underlying patterns in the data, leading to less reliable results.

To address this limitation, we propose introducing a Monte Carlo k-folding technique in our future work. Unlike the current method, this technique systematically partitions the entire dataset into k subsets (or "folds"). In each iteration, one of these folds is used as the test set, while the remaining k-1 folds are used for training the model. This process is repeated k times, with each fold serving as the test set exactly once.

However, instead of just performing standard k-fold cross-validation, we will integrate Monte Carlo sampling into the process. This means that within each fold, we will randomly sample points multiple times across different iterations, ensuring that over the course of training, every point in the dataset is eventually selected and tested multiple times, rather than the current implementation where it can not be ensured that every point is selected at least once. This repeated random sampling within the k-fold framework allows us to estimate the expected value of the loss more accurately for each point.

By employing Monte Carlo k-folding, we can better assess the model's performance across the entire dataset, rather than relying on a small, potentially unrepresentative subset of points. This approach will lead to a more thorough evaluation, ensuring that our model is tested on a more diverse and representative sample of the data, ultimately resulting in a more accurate and generalized model. This refinement is expected to yield a more robust understanding of the model's predictive capabilities and improve its reliability in practical applications.

8.2 Implementing Noise

We currently calculate our known β and γ values by simulating the QAOA quantum circuits on a Qiskit simulator without noise. While this simulation is efficient, it lacks accuracy in representing the noisy testing environment of the RPI-IBM Quantum System One.

To address this issue, we propose calculating the known β and γ by simulating the QAOA quantum circuits on a Qiskit simulator with noise. Although this approach does not perfectly replicate the noisy testing environment of the RPI-IBM Quantum System One, it would provide a more accurate and realistic simulation when classically calculating β and γ values.

9 References

[1] Hidary, Jack. 2019. Quantum Computing : An Applied Approach. Springer Nature.

- 361 [2] Jain, Nishant, Brian Coyle, Elham Kashefi, and Niraj Kumar. 2022. “Graph Neural
362 Network Initialisation of Quantum Approximate Optimisation.” *Quantum* 6 (November): 861.
363 <https://doi.org/10.22331/q-2022-11-17-861>.
- 364 [3] Jin, Allison, and Xiao-Yang Liu. 2023. “A Fast Machine Learning Algorithm for the MaxCut
365 Problem.” 2023 IEEE151 MIT Undergraduate Research Technology Conference (URTC), October,
366 1–5. <https://doi.org/10.1109/urtc60662.2023.10534996>.
- 367 [4] Liang, Zhiding, Gang Liu, Zheyuan Liu, Jinglei Cheng, Tianyi Hao, Kecheng Liu, Hang Ren, et al.
368 n.d. “Graph Learning for Parameter Prediction of Quantum Approximate Optimization Algorithm.”
369 Accessed August 18, 2024. <https://arxiv.org/pdf/2403.03310>.
- 370 [5] Lu, Yicheng, and Xiao- Yang Liu. 2023. “Reinforcement Learning for Ising Model.” *NeurIPS* .
- 371 [6] Lucas, Andrew. 2014. “Ising Formulations of Many NP Problems.” *Frontiers in Physics* 2.
372 <https://doi.org/10.3389/fphy.2014.00005>.
- 373 [7] Google Quantum AI, "Quantum Approximate Optimization Algorithm for the Ising Model," 2024.
374 Available: https://quantumai.google/cirq/experiments/qaoa/qaoa_ising.
- 375 [8] Uzair Aslam Bhatti, Hao Tang, Guilu Wu, Shah Marjan, and Aamir Hussain. 2023. “Deep Learn-
376 ing with Graph Convolutional Networks: An Overview and Latest Applications in Computational
377 Intelligence” 2023 (February): 1–28. <https://doi.org/10.1155/2023/8342104>.
- 378 [9] Zhang, Si, Hanghang Tong, Jiejun Xu, and Ross Maciejewski. 2019. “Graph Con-
379 volutional Networks: A Comprehensive Review.” *Computational Social Networks* 6 (1).
380 <https://doi.org/10.1186/s40649-019-0069-y>.